

ML LAB 14

Name: Mrigank Jain

Semester: 5

Section: F

SRN: PES2UG23CS352

Introduction:

The objective of this laboratory exercise was to design, implement, and train a Convolutional Neural Network (CNN) using the PyTorch framework. The specific task was to build a model capable of image classification , accurately distinguishing between three distinct hand gestures: 'rock', 'paper', and 'scissors'. The project involved completing a boilerplate notebook, preprocessing the "Rock Paper Scissors" dataset, defining a custom CNN architecture, and training the model to evaluate its final performance on unseen test data.

Model Architecture:

The classification task was handled by a custom CNN model named RPS_CNN. The network consists of a feature extraction block (convolutional layers) and a classifier block (fully-connected layers).

Convolutional Block (self.conv_block)

The feature extractor is composed of three sequential convolutional layers, each followed by a ReLU activation and a MaxPool2d(2) layer. The input images were preprocessed to 128 x 128 pixels, and the spatial dimensions are reduced by a factor of 2 in each MaxPool layer.

- **Block 1:** Conv2d (3 input channels, 16 output channels, kernel_size=3, padding=1), ReLU, MaxPool2d(2).
- **Block 2:** Conv2d (16 input channels, 32 output channels, kernel_size=3, padding=1), ReLU, MaxPool2d(2).
- **Block 3:** Conv2d (32 input channels, 64 output channels, kernel_size=3, padding=1), ReLU, MaxPool2d(2).

Fully-Connected Classifier (self.fc)

The output of the final convolutional block is passed to the classifier. After the three MaxPool(2) layers, the spatial size is reduced to 16 x 16, resulting in 64 feature maps.

- **Flatten Layer:** The feature maps are flattened, yielding an input size of 16384 features ($64 \times 16 \times 16$).
- **Layer 1:** A linear layer with weights mapping 16384 inputs to 256 outputs, followed by a ReLU activation.
- **Regularization:** A Dropout layer with a probability $p=0.3$ was applied to prevent overfitting.

- **Output Layer:** The final linear layer maps 256 inputs to 3 outputs, corresponding to the three classes.

Training and Performance:

Hyperparameters

The model was trained using the following key hyperparameters:

- **Optimizer:** Adam
- **Loss Function (Criterion):** CrossEntropyLoss.
- **Learning Rate (lr):** 0.001.
- **Number of Epochs:** 10.
- **Batch Size:** 32.

Test Accuracy

The model was evaluated on the 20% unseen test set.

The final test accuracy is 99.54%

Conclusion and Analysis:

The CNN model demonstrated excellent performance on the image classification task, achieving a final test accuracy of 99.54%. The architecture, which utilized three convolutional blocks for feature extraction and a fully-connected layer with Dropout for classification, proved highly effective in learning the distinct features of the hand gestures.

Challenges:

A key challenge encountered during the initial setup was ensuring the data loading pipeline was correctly configured, specifically by correcting the hardcoded source path (src\root) to match the actual download location returned by kagglehub.

Suggested Improvements:

1. **Data Augmentation:** To improve the model's generalization capabilities, the input transforms should be expanded to include more robust data augmentation techniques such as random rotation, scaling, and shear transformations.

2. **Transfer Learning:** A substantial performance increase could likely be achieved by using a state-of-the-art pre-trained model as a feature extractor. This technique would allow the model to leverage features learned from millions of general images, providing a much stronger base for classification.